

Importing libraries

In [16]:

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
# for the Q-Q plots
#import scipy.stats as stats
%matplotlib inline
import pandas as pd
pd.options.display.float_format = '{:.2f}'.format
#from pandas.io.json import json_normalize
```

Loading dataset for receipts

In [17]:

```
receipts = pd.read_excel("receipts.xlsx")
```

In [18]:

```
receipts.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1119 entries, 0 to 1118
Data columns (total 15 columns):
 _id/$oid                1119 non-null object
 bonusPointsEarned       544 non-null float64
 bonusPointsEarnedReason 544 non-null object
 createDate              1119 non-null int64
 dateScanned             1119 non-null int64
 finishedDate            568 non-null float64
 modifyDate              1119 non-null int64
 pointsAwardedDate       537 non-null float64
 pointsEarned            609 non-null float64
 purchaseDate            671 non-null float64
 purchasedItemCount      635 non-null float64
 rewardsReceiptItemList  679 non-null object
 rewardsReceiptStatus    1119 non-null object
 totalSpent              684 non-null float64
 userId                  1119 non-null object
dtypes: float64(7), int64(3), object(5)
memory usage: 131.3+ KB
```

In [19]:

```
receipts.head()
```

Out[19]:

	_id/\$oid	bonusPointsEarned	bonusPointsEarnedReason	createDate	d
0	5ff1e1eb0a720f0523000575	500.00	Receipt number 2 completed, bonus point schedu...	1609687531000	160
1	5ff1e1bb0a720f052300056b	150.00	Receipt number 5 completed, bonus point schedu...	1609687483000	160
2	5ff1e1f10a720f052300057a	5.00	All-receipts receipt bonus	1609687537000	160
3	5ff1e1ee0a7214ada100056f	5.00	All-receipts receipt bonus	1609687534000	160
4	5ff1e1d20a7214ada1000561	5.00	All-receipts receipt bonus	1609687506000	160



Identifying numerical and categorical variables

In [20]:

```
# make lists of variable types

temporal = [var for var in receipts.columns if 'date' in var or 'Date' in var]

discrete = [
    var for var in receipts.columns if receipts[var].dtype != 'O'
    and len(receipts[var].unique()) < 20 and var not in temporal
]

continuous = [
    var for var in receipts.columns if receipts[var].dtype != 'O'
    if var not in discrete and var != '_id'
    and var not in temporal
]

categorical = [var for var in receipts.columns if receipts[var].dtype == 'O'
                and var not in temporal and var not in discrete]

print(f'There are {len(continuous)} continuous variables')
print(f'There are {len(discrete)} discrete variables')
print(f'There are {len(temporal)} temporal variables')
print(f'There are {len(categorical)} categorical variables')
```

There are 3 continuous variables
There are 1 discrete variables
There are 6 temporal variables
There are 5 categorical variables

In [21]:

```
continuous
```

Out[21]:

```
['pointsEarned', 'purchasedItemCount', 'totalSpent']
```

In [22]:

```
discrete
```

Out[22]:

```
['bonusPointsEarned']
```

In [23]:

temporal

Out[23]:

```
[ 'createDate',
  'dateScanned',
  'finishedDate',
  'modifyDate',
  'pointsAwardedDate',
  'purchaseDate']
```

In [24]:

categorical

Out[24]:

```
[ '_id/$oid',
  'bonusPointsEarnedReason',
  'rewardsReceiptItemList',
  'rewardsReceiptStatus',
  'userId']
```

Quantifying missing data

In [30]:

receipts.isnull().sum()

Out[30]:

_id/\$oid	0
bonusPointsEarned	575
bonusPointsEarnedReason	575
createDate	0
dateScanned	0
finishedDate	551
modifyDate	0
pointsAwardedDate	582
pointsEarned	510
purchaseDate	448
purchasedItemCount	484
rewardsReceiptItemList	440
rewardsReceiptStatus	0
totalSpent	435
userId	0
dtype: int64	

percentage of missing values in variables

In [32]:

```
# alternatively, we can use the mean() method after isnull() to visualise the percentage of missing values for each variable
percentage_null_values= receipts.isnull().mean()
for key,value in percentage_null_values.items():
    if value >0:
        print(key, ":", value*100)
```

```
bonusPointsEarned : 51.385165326184094
bonusPointsEarnedReason : 51.385165326184094
finishedDate : 49.240393208221626
pointsAwardedDate : 52.01072386058981
pointsEarned : 45.57640750670242
purchaseDate : 40.03574620196604
purchasedItemCount : 43.25290437890974
rewardsReceiptItemList : 39.32082216264522
totalSpent : 38.8739946380697
```

A considerable fraction of values are missing from the above mentioned variables. Missing values for certain variables are a major concern:

1. finishedDate- for 49%(almost half) of the receipts we don't know when do they become invalid(assuming that the date on which a receipt finishes processing is the date on which it becomes invalid)

In [34]:

```
receipts["pointsEarned"].unique()
```

Out[34]:

```
array([5.00000e+02, 1.50000e+02, 5.00000e+00, 7.50000e+02, 2.50000e+02,
       1.00000e+02, 8.85000e+03, 3.00000e+02, nan, 3.89200e+02,
       1.85000e+02, 3.50000e+01, 6.50000e+02, 5.50000e+01, 5.00000e+01,
       3.55000e+02, 6.00000e+02, 1.75000e+03, 3.50000e+02, 2.25000e+02,
       2.75000e+02, 2.50000e+01, 7.55000e+02, 1.80000e+03, 8.10000e+02,
       3.05000e+02, 9.44980e+03, 9.12000e+01, 8.25000e+02, 3.50600e+02,
       1.25000e+02, 7.93100e+02, 2.00000e+02, 3.25000e+03, 0.00000e+00,
       4.00500e+03, 2.00500e+03, 8.41200e+02, 5.75000e+03, 3.75000e+03,
       8.70000e+03, 7.60000e+02, 7.80000e+02, 9.20000e+03, 1.00500e+03,
       1.99960e+03, 1.89200e+02, 8.95000e+03, 8.85000e+02, 8.00000e+02,
       2.95000e+02, 6.82400e+02, 8.37400e+02, 2.37800e+02, 1.60000e+02,
       8.55700e+02, 6.05700e+02, 2.41670e+03, 1.80640e+03, 4.05700e+02,
       1.51690e+03, 1.65830e+03, 2.68580e+03, 8.79100e+02, 3.65940e+03,
       9.34400e+02, 8.77700e+02, 9.22100e+02, 1.54180e+03, 1.00000e+03,
       5.74400e+02, 5.06000e+01, 2.05550e+03, 5.85000e+03, 4.85000e+03,
       1.73600e+02, 9.85000e+03, 5.09000e+01, 2.30000e+03, 6.73000e+02,
       4.05900e+02, 2.14330e+03, 1.55000e+03, 9.86500e+02, 5.83400e+02,
       4.48050e+03, 3.37990e+03, 3.23600e+02, 6.25730e+03, 2.49770e+03,
       1.17870e+03, 1.44700e+03, 1.72950e+03, 1.47620e+03, 1.70800e+03,
       1.07750e+03, 1.13510e+03, 1.04430e+03, 6.40700e+02, 5.23600e+02,
       4.87700e+02, 7.13720e+03, 1.20500e+03, 4.00000e+02, 1.01998e+04,
       8.50000e+02, 9.46000e+01, 1.85000e+03, 8.40000e+02, 1.04980e+03,
       9.40000e+02, 4.94470e+03, 3.75000e+02, 1.49950e+03, 2.09900e+02,
       2.09800e+02, 2.10000e+02, 2.09500e+02, 7.89200e+02, 3.50000e+03])
```

1. pointsEarned- 45% of the values for the 'pointsEarned' field are missing. If we look at the unique values for 'pointsEarned', we do not have a zero value. This means that points were earned for certain receipts but the data was not captured and that's why the large number of missing values.
1. purchasedItemCount- large number of missing values will pose problems for deciding if users who bought more than one unit of a product qualify for special offers/bonus points that require them to purchase certain amount of particular products/brands.
1. totalSpent, rewardsReceiptItemList- Since data for these two fields is missing, it is natural that we don't have information about points earned(pointsEarned field) for those transactions.

Examining percentage of different category values for categorical variables

Here, the categorical variable of my interest is bonusPointsEarnedReason

In [37]:

```
freq_reasons = 100*(receipts['bonusPointsEarnedReason'].value_counts() / len(receipts))
print(freq_reasons.map('{:,.2f} %'.format))
```

All-receipts receipt bonus

16.35 %

Receipt number 1 completed, bonus point schedule DEFAULT (5cefdcacf3693e0b50e83a36) 10.63 %

COMPLETE_NONPARTNER_RECEIPT

6.34 %

COMPLETE_PARTNER_RECEIPT

3.49 %

Receipt number 3 completed, bonus point schedule DEFAULT (5cefdcacf3693e0b50e83a36) 2.77 %

Receipt number 2 completed, bonus point schedule DEFAULT (5cefdcacf3693e0b50e83a36) 2.68 %

Receipt number 5 completed, bonus point schedule DEFAULT (5cefdcacf3693e0b50e83a36) 2.41 %

Receipt number 4 completed, bonus point schedule DEFAULT (5cefdcacf3693e0b50e83a36) 2.32 %

Receipt number 6 completed, bonus point schedule DEFAULT (5cefdcacf3693e0b50e83a36) 1.61 %

Name: bonusPointsEarnedReason, dtype: object


In []:

Plotting outliers using boxplot

In [45]:

```
fig, axs = plt.subplots(ncols=3, nrows=1, figsize=(15,5))
axs = axs.flatten()

for i, var in enumerate(continuous):
    sns.boxplot(receipts[receipts[var].notnull()][var], ax=axs[i], orient='h');
```




By looking at the boxplots we can say that there is a significantly large range of outliers.

Examining distributions of continuous variables

In [46]:

```
fig, axs = plt.subplots(ncols=3, nrows=1, figsize=(15,5))
axs = axs.flatten()

for i, var in enumerate(continuous):
    sns.distplot(receipts[receipts[var].notnull()][var], ax=axs[i]);
```



The distributions for all the three continuous variables are right-skewed.

Both, the boxplots and value distribution diagrams suggest that there are a lot of outliers, but we can't say anything about the significance of these outliers at this point.


Examining correlation between variables

In [48]:

```
import seaborn as sns
f, ax = plt.subplots(figsize=(10, 10))
corr = receipts.corr()
sns.heatmap(corr, mask=np.zeros_like(corr, dtype=np.bool), square=True, ax=ax
)
```

Out[48]:

<matplotlib.axes._subplots.AxesSubplot at 0x1501ba17d88>



No significant correlations found.

In []: